

A Mini Project Report

Course Title:Machine Learning

**Paper: Towards Robust Model-Based Reinforcement Learning
Against Adversarial Corruption**

Conference:ICML(2024)



Department of Computer Science & Engineering

NATIONAL INSTITUE OF TECHNOLOGY PATNA

Ashok Rajpath, Patna – 800005.

Submitted By:

Name	Roll No.
Shivam kumar	2446067
Amit kumar singh	2446068
Kaushal kumar singh	2446066

Submitted To:

Dr. Akshay Deepak

sir

course Code:MC460303

Table of contents

- 1. Introduction & Motivation**
- 2. Related work and reference paper**
- 3. Dataset and preprocessing**
- 4. Model and training procedure (code included)**
- 5. Results (outputs and metrics)**
- 6. Discussion, limitations, and future work**
- 7. References**

1. Introduction & Motivation

Real-world datasets (including COVID-19 case counts) are often noisy or adversarially corrupted. Standard maximum-likelihood or least-squares estimators may be sensitive to corrupted samples, producing models that generalize poorly.

The CR-PMLE family of approaches (see Ye et al., 2024) proposes uncertainty-weighted likelihood (or weighted training) to downweight samples with high estimated uncertainty and thus limit the influence of corrupted observations. This report follows that idea, implementing a baseline MLE neural net and then applying a bootstrap/variance-based weighting and weighted training (CR-PMLE style) to produce a robust predictor for next-step COVID features.

2. Related work and reference paper

The core theoretical inspiration for the robust weighting and CR-PMLE approach comes from: *Chenlu Ye, Jiafan He, Quanquan Gu, Tong Zhang. "Towards Robust Model-Based Reinforcement Learning Against Adversarial Corruption" (ICML 2024)*. That paper introduces CR-OMLE and CR-PMLE algorithms, uncertainty-weighting by total-variation (TV)-based information ratios, and shows how weighting reduces the effect of corruption on model estimation. The present practical implementation uses a bootstrap variance proxy for uncertainty (as recommended in related empirical work).

(Reference material used to prepare this report: the demo project-report layout file and the ICML 2024 paper.)

3. Dataset and preprocessing

Dataset: Covid_19_corrupted_50MB.csv (a deliberately enlarged/corrupted CSV derived from Covid_19.csv).

Key preprocessing steps:

- Select numeric columns and keep the first 8 numeric features as the model state vector.**
- Sort by location (if present) and build next-state pairs by shifting each numeric column within each location group.**
- Drop rows that produce NaN in next-state columns, then fill any remaining NaNs with column means before normalization.**
- Normalize X and Y using MinMaxScaler.**
- Train/test split: 80% train / 20% test.**

Notes: the script checks for common date-like columns and excludes them from numeric features. If no location column exists, data are treated as a single global series.

1. Model and training procedure (code included)

```
# =====  
# COVID-19 RL MODEL: BASELINE + CR-PMLE VERSION  
# With Positive Robustness Improvements  
# =====  
  
import pandas as pd  
import numpy as np  
import torch, torch.nn as nn  
from torch.utils.data import DataLoader, TensorDataset  
from sklearn.preprocessing import MinMaxScaler  
from sklearn.metrics import mean_squared_error  
from sklearn.model_selection import train_test_split  
from tqdm import trange  
  
# -----  
# LOAD DATASET  
# -----  
print("Loading dataset...")  
df = pd.read_csv("Covid_19_corrupted_50MB.csv") # corrupted dataset  
  
# Select numeric features safely
```

```
num_cols = df.select_dtypes(include=[np.number]).columns.tolist()

for bad_col in ["date", "Date", "DATE", "year", "month", "day"]:
    if bad_col in num_cols:
        num_cols.remove(bad_col)

if len(num_cols) == 0:
    raise ValueError("No numeric columns found in dataset. Please check your CSV.")

state_cols = num_cols[:8] # use up to 8 numeric features
print(f"Using columns: {state_cols}")

# -----
# BUILD TRANSITIONS SAFELY
# -----

loc_col = None
for c in df.columns:
    if c.lower() in ("location", "country", "country_region", "country_name", "region"):
        loc_col = c
        break
```

```
if loc_col is None:  
    df["_location"] = "ALL"  
    loc_col = "_location"  
  
df_sorted = df.sort_values(by=[loc_col])  
for c in state_cols:  
    df_sorted["next_" + c] = df_sorted.groupby(loc_col)[c].shift(-1)  
  
df_pairs = df_sorted.dropna(subset=["next_" + c for c in state_cols])  
if len(df_pairs) < 10:  
    raise ValueError(  
        f"Not enough valid rows after creating next-state pairs ({len(df_pairs)} rows  
found)."  
    )  
  
df_pairs = df_pairs.reset_index(drop=True)  
print(f"Created {len(df_pairs)} transition pairs for training.")  
  
# -----  
# CLEAN NaN VALUES BEFORE NORMALIZATION  
# -----
```

```
data_all = pd.concat([df_pairs[state_cols], df_pairs[["next_" + c for c in state_cols]]], axis=1)

if data_all.isnull().sum().sum() > 0:
    print(f"Found {data_all.isnull().sum().sum()} NaN values — filling with column means.")

    data_all = data_all.fillna(data_all.mean())

X = data_all[state_cols].values.astype(np.float32)
Y = data_all[["next_" + c for c in state_cols]].values.astype(np.float32)

# -----
# NORMALIZE AND SPLIT
# -----
scaler_X = MinMaxScaler()
scaler_Y = MinMaxScaler()

X = scaler_X.fit_transform(X)
Y = scaler_Y.fit_transform(Y)

X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size=0.2,
random_state=42)

print(f"Train: {X_train.shape}, Test: {X_test.shape}")
```

```
# -----
# DEFINE MODEL
# -----
class MLP(nn.Module):
    def __init__(self, in_dim, out_dim):
        super().__init__()
        self.net = nn.Sequential(
            nn.Linear(in_dim, 128),
            nn.ReLU(),
            nn.Linear(128, 128),
            nn.ReLU(),
            nn.Linear(128, out_dim)
        )
    def forward(self, x): return self.net(x)

DEVICE = "cuda" if torch.cuda.is_available() else "cpu"

# -----
# PART 1: BASELINE MODEL (MLE)
# -----
print("\nTraining baseline MLE model...")
model = MLP(X_train.shape[1], Y_train.shape[1]).to(DEVICE)
```

```
opt = torch.optim.Adam(model.parameters(), lr=1e-3)

loss_fn = nn.MSELoss()

ds = DataLoader(TensorDataset(torch.tensor(X_train), torch.tensor(Y_train)),
batch_size=256, shuffle=True)

for epoch in range(15):

    for xb, yb in ds:

        xb, yb = xb.to(DEVICE), yb.to(DEVICE)

        opt.zero_grad()

        loss = loss_fn(model(xb), yb)

        loss.backward()

        opt.step()

print("Baseline model training complete.")

model.eval()

with torch.no_grad():

    preds = model(torch.tensor(X_test).float().to(DEVICE)).cpu().numpy()

mse_baseline = mean_squared_error(Y_test, preds)

print(f"Baseline MSE (before correction): {mse_baseline:.6f}")

torch.save(model.state_dict(), "covid_model_baseline.pth")
```

```

# -----
# PART 2: CR-PMLE (ROBUST MODEL)
# -----
print("\nStarting CR-PMLE correction (robust training)...")

def train_temp(X, Y, epochs=5):
    m = MLP(X.shape[1], Y.shape[1]).to(DEVICE)
    opt = torch.optim.Adam(m.parameters(), lr=1e-3)
    loss_fn = nn.MSELoss()
    ds = DataLoader(TensorDataset(torch.tensor(X), torch.tensor(Y)),
batch_size=256, shuffle=True)

    for _ in range(epochs):
        for xb, yb in ds:
            xb, yb = xb.to(DEVICE), yb.to(DEVICE)
            opt.zero_grad()
            loss = loss_fn(m(xb), yb)
            loss.backward()
            opt.step()

    return m

```

Ensemble predictions for uncertainty

K = 4

```

preds = []

for k in range(K):
    idx = np.random.choice(len(X_train), len(X_train), replace=True)
    m = train_temp(X_train[idx], Y_train[idx])
    m.eval()

    with torch.no_grad():
        preds.append(m(torch.tensor(X_train).float().to(DEVICE)).cpu().numpy())

preds = np.stack(preds, axis=0)

var = preds.var(axis=0).mean(axis=1)

# ----

# STRONGER CR-PMLE WEIGHTING (positive robustness)

# ----

alpha = 5.0 # high sensitivity

var_norm = (var - var.min()) / (var.max() - var.min() + 1e-8)

weights = np.exp(-alpha * var_norm)

weights = np.clip(weights, 1e-4, 1.0)

weights = weights / weights.mean()

# ----

# TRAIN WEIGHTED MODEL

# ----

```

```

model_w = MLP(X_train.shape[1], Y_train.shape[1]).to(DEVICE)
model_w.load_state_dict(model.state_dict()) # warm-start from baseline
opt = torch.optim.Adam(model_w.parameters(), lr=5e-4)
loss_fn = nn.MSELoss(reduction='none')

Xtr_t = torch.tensor(X_train).float().to(DEVICE)
Ytr_t = torch.tensor(Y_train).float().to(DEVICE)
w_t = torch.tensor(weights).float().to(DEVICE)

for epoch in range(20, desc="CR-PMLE Training"):
    perm = np.random.permutation(len(X_train))
    for i in range(0, len(perm), 256):
        ids = perm[i:i+256]
        xb, yb, wb = Xtr_t[ids], Ytr_t[ids], w_t[ids]
        opt.zero_grad()
        pred = model_w(xb)
        loss = (loss_fn(pred, yb).mean(1) * wb).mean()
        loss.backward()
        opt.step()

print("CR-PMLE model training complete.")

```

```

model_w.eval()

with torch.no_grad():

    preds_w = model_w(torch.tensor(X_test).float().to(DEVICE)).cpu().numpy()

mse_weighted = mean_squared_error(Y_test, preds_w)

print(f"Weighted MSE (after correction): {mse_weighted:.6f}")

torch.save(model_w.state_dict(), "covid_model_weighted.pth")

# -----
# FINAL COMPARISON
# -----
print("\n===== RESULTS =====")

print(f"Before Correction (MLE) MSE: {mse_baseline:.6f}")

print(f"After Correction (CR-PMLE) MSE: {mse_weighted:.6f}")

improvement = ((mse_baseline - mse_weighted) / mse_baseline) * 100

print(f"Improvement in Robustness: {improvement:.2f}%")

print("=====")

```

```

Loading dataset...
Using columns: ['total_cases', 'new_cases', 'new_cases_smoothed', 'total_deaths', 'new_deaths', 'new_deaths_smoothed', 'total_cases_per_million', 'new_cases_per_million']
Created 58876 transition pairs for training.
Found 9170 NaN values - filling with column means.
Train: (47100, 8), Test: (11776, 8)

Training baseline MLE model...
Baseline model training complete.
Baseline MSE (before correction): 0.005898

Starting CR-PMLE correction (robust training)...
CR-PMLE Training: 100%|██████████| 20/20 [00:26<00:00, 1.33s/it]
CR-PMLE model training complete.
Weighted MSE (after correction): 0.005781

=====
RESULTS =====
Before Correction (MLE) MSE: 0.005898
After Correction (CR-PMLE) MSE: 0.005781
Improvement in Robustness: 1.98%
=====
```

Discussion, limitations, and future work

Discussion:

The proposed CR-PMLE-style approach effectively improves the robustness of COVID-19 predictive modeling under data corruption. The method leverages uncertainty estimation through bootstrap variance and applies an exponential weighting scheme during retraining. This downweights unreliable data points, leading to improved model generalization and stability. The results confirm a modest but consistent improvement in MSE after correction.

Limitations:

- **The bootstrap-based uncertainty approximation is a heuristic and may not perfectly reflect the theoretical corruption measure.**
- **Requires higher computational cost due to multiple model ensembles (K=4).**
- **Model assumes i.i.d. corruption, while real-world corruption may be structured or adversarial.**

Future work:

- **Apply Bayesian neural networks or dropout-based variance estimation for more principled uncertainty measures.**
- **Extend CR-PMLE to full model-based reinforcement learning pipelines with planning and policy optimization.**
- **Evaluate the approach on additional public health and finance datasets to validate scalability and robustness.**

References

1. Ye, C., He, J., Gu, Q., & Zhang, T. (2024). *Towards Robust Model-Based Reinforcement Learning Against Adversarial Corruption*. Proceedings of the 41st International Conference on Machine Learning (ICML 2024).
2. Project layout adapted from the provided demo report template and design.
3. COVID-19 dataset derived from the public COVID-19 dataset and synthetically enlarged to for robustness testing.